

Logistic regression

September 28, 2021

1 Logistic regression

The logistic regression is very useful when the outcome (dependant variable) is discrete, or so said categorical random variable. It is usee to describe 2 outcomes. the formula of logistic regression is:

$$\ln\left(\frac{p}{1-p}\right) = m + kx$$

where p is the probability of winning outcome, or outcome is 1, and $\frac{p}{1-p}$ is odds. x is the independant variable. m and k are coefficients. \ln is the logrithmic operation.

The formula can be changed to another form:


$$p(y = 1) = \frac{1}{1 + e^{-(m+kx)}}$$

It is a sigmoid function (**S-Curve**). With diffrent x value, the probability of outcome being 1 is like the following plot (p):

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
plt.style.use('seaborn')
plt.rcParams.update({'figure.figsize': (12,6),\
                        'figure.titlesize':16,\
                        'axes.titlesize':16,\
                        'axes.labelsize':14,\
                        'xtick.labelsize':14,\
                        'ytick.labelsize':14,\
                        })


[2]: m = -5
k = 1
x = np.linspace(0,10)
p_func = lambda x: 1/(1+np.exp(-(k*x+m)))
fig, ax = plt.subplots()
ax.plot(x,p_func(x),label='logistic', color='coral')
ax.plot([0,abs(m)], [0.5,0.5], dashes=(4,4), color='grey')
ax.plot([abs(m),abs(m)], [-.1,.5], dashes=(4,4), color='grey')
ax.set_title('Sigmoid Function')
```

```
ax.set_xlabel('x')
ax.set_ylabel('p')
plt.show()
```



When k increase to large number, it quickly becomes a sharp change function, which very similar to categorial (2 values)

```
[3]: m = -50
k = 10
x = np.linspace(0,10)
p_func = lambda x: 1/(1+np.exp(-(k*x+m)))
fig, ax = plt.subplots()
ax.plot(x,p_func(x),label='logistic', color='coral')
ax.set_title('Sigmoid Function')
ax.set_xlabel('x')
ax.set_ylabel('p')
plt.show()
```



A multi variate sigmoid function would be:

$$p(y = 1) = \frac{1}{1 + e^{-m + \sum_i k_i x_i}}$$

1.1 Single Variable Logistic Regression

the exam dataset This dataset is from the book “Mastering Python Data Analysis”, by Luiz Felipe Martins.

```
[4]: studytime = [0, 0, 1.5, 2, 2.5, 3, 3.5, 4, 4, 4, 4, 5.5, 6, 6.5, 7, 7, 8.5, 9, 9, 9,
    ↪9, 10.5, 10.5, 12, 12, 12, 12.5, 13, 14, 15, 16, 18]
passed = [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 1, 1,
    ↪0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1]
data = pd.DataFrame(data=np.array([studytime, passed]).T, columns=['Study_
    ↪Time', 'Passed'])
data
```

```
[4]:
```


	Study Time	Passed
0	0.0	0.0
1	0.0	0.0
2	1.5	0.0
3	2.0	0.0
4	2.5	0.0
5	3.0	0.0
6	3.5	0.0
7	4.0	0.0

8	4.0	0.0
9	4.0	0.0
10	5.5	0.0
11	6.0	1.0
12	6.5	0.0
13	7.0	1.0
14	7.0	1.0
15	8.5	0.0
16	9.0	1.0
17	9.0	1.0
18	9.0	0.0
19	10.5	1.0
20	10.5	1.0
21	12.0	1.0
22	12.0	1.0
23	12.0	1.0
24	12.5	1.0
25	13.0	1.0
26	14.0	1.0
27	15.0	1.0
28	16.0	1.0
29	18.0	1.0


```
[5]: data.dtypes
```

```
[5]: Study Time    float64  
     Passed       float64  
     dtype: object
```

```
[6]: ax = data['Study Time'].hist(bins=18, ec='white')  
     ax.set_xlabel('Study Time')  
     ax.set_ylabel('# Students')  
     plt.show()
```



```
[7]: ax = data.plot(x='Study Time', y='Passed', kind='scatter')
```



We can see that from study time 5 to 10, there are some student passed the exam, some not. but above 10, all students passed. below 5, all student not pass.

use scikit-learn


```
[8]: from sklearn.linear_model import LogisticRegression
      clf = LogisticRegression()
```

```
clf.fit(data[['Study Time']], data['Passed'])
print(clf.coef_)
print(clf.intercept_)
```

```
[[0.74291088]]
```

```
[-5.37187834]
```

```
[9]: fig, ax = plt.subplots()
test_x = np.linspace(min(studytime)-0.5, max(studytime)+0.5,100)
ax.plot(test_x, clf.predict(test_x.reshape(-1,1)), label='scikit learn',
        color='coral')
ax.set_xlabel('Study Time')
ax.set_ylabel('p')
plt.show()
```



Scikit package logistic regression prediction function, will gives a strict classifier result, strictly separated with x around 7.5.

Use statsmodels

```
[10]: import statsmodels.api as sm

st_train_x = sm.add_constant(data[['Study Time']])
log_reg = sm.Logit(data['Passed'], st_train_x)
log_result = log_reg.fit()
log_result.summary()
```

Optimization terminated successfully.

Current function value: 0.251107

Iterations 8

C:\Python39\lib\site-packages\statsmodels\tsa\tsatools.py:142: FutureWarning: In a future version of pandas all arguments of concat except for the argument 'objs' will be keyword-only


```
x = pd.concat(x[:, :order], 1)
```

```
[10]: <class 'statsmodels.iolib.summary.Summary'>
```

```
"""
                                Logit Regression Results
=====
Dep. Variable:                Passed    No. Observations:                30
Model:                        Logit     Df Residuals:                    28
Method:                        MLE      Df Model:                        1
Date:                          Tue, 28 Sep 2021    Pseudo R-squ.:                0.6366
Time:                          09:59:55    Log-Likelihood:               -7.5332
converged:                      True      LL-Null:                       -20.728
Covariance Type:                nonrobust    LLR p-value:                  2.791e-07
=====
               coef      std err          z      P>|z|      [0.025      0.975]
-----
const          -5.7980      2.240      -2.588      0.010     -10.188      -1.408
Study Time       0.8020      0.297       2.703      0.007       0.220       1.384
=====
"""
```

The coef for const and x, are different from scikit package.

```
[11]: y_predict = log_result.predict(sm.add_constant(test_x.reshape(-1,1)))
fig, ax = plt.subplots()
ax.plot(test_x, y_predict, label='statsmodels', color='steelblue')
ax.set_xlabel('Study Time')
ax.set_ylabel('p')
plt.show()
```



The prediction is not a sharp change, but is a **smooth S-Curve**. So it is not like a classifier, still a continuous regression result. Between 5 ~ 10 is the transition period also fits for our observation from the train data.

Scikit-learn package predict result seems give a better classifier result, the results are only 2.

1.2 Multi Variables Logistic Regression

the heart disease dataset Download this dataset from Kaggle website, csv format file is available.

1. Demographic:
 - Sex: male or female(Nominal)
 - Age: Age of the patient;(Continuous - Although the recorded ages have been truncated to whole numbers, the concept of age is continuous)
2. Behavioral
 - Current Smoker: whether or not the patient is a current smoker (Nominal)
 - Cigs Per Day: the number of cigarettes that the person smoked on average in one day.(can be considered continuous as one can have any number of cigarettes, even half a cigarette.)
3. Medical(history)
 - BP Meds: whether or not the patient was on blood pressure medication (Nominal)
 - Prevalent Stroke: whether or not the patient had previously had a stroke (Nominal)
 - Prevalent Hyp: whether or not the patient was hypertensive (Nominal)

- Diabetes: whether or not the patient had diabetes (Nominal)
4. Medical(current)
- Tot Chol: total cholesterol level (Continuous)
 - Sys BP: systolic blood pressure (Continuous)
 - Dia BP: diastolic blood pressure (Continuous)
 - BMI: Body Mass Index (Continuous)
 - Heart Rate: heart rate (Continuous - In medical research, variables such as heart rate though in fact discrete, yet are considered continuous because of large number of possible values.)
 - Glucose: glucose level (Continuous)
5. Predict variable (desired target)
- 10 year risk of coronary heart disease CHD (binary: “1”, means “Yes”, “0” means “No”)

```
[12]: import pathlib
```

```
file_path = pathlib.Path('D:/Edu/data_resource/framingham.csv')
df = pd.read_csv(file_path,header=0)
df
```

```
[12]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	\	
0	1	39	4.0	0	0.0	0.0		
1	0	46	2.0	0	0.0	0.0		
2	1	48	1.0	1	20.0	0.0		
3	0	61	3.0	1	30.0	0.0		
4	0	46	3.0	1	23.0	0.0		
...		
4233	1	50	1.0	1	1.0	0.0		
4234	1	51	3.0	1	43.0	0.0		
4235	0	48	2.0	1	20.0	NaN		
4236	0	44	1.0	1	15.0	0.0		
4237	0	52	2.0	0	0.0	0.0		
...		
	prevalentStroke	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	\
0	0	0	0	195.0	106.0	70.0	26.97	
1	0	0	0	250.0	121.0	81.0	28.73	
2	0	0	0	245.0	127.5	80.0	25.34	
3	0	1	0	225.0	150.0	95.0	28.58	
4	0	0	0	285.0	130.0	84.0	23.10	
...	

4233	0	1	0	313.0	179.0	92.0	25.97
4234	0	0	0	207.0	126.5	80.0	19.71
4235	0	0	0	248.0	131.0	72.0	22.00
4236	0	0	0	210.0	126.5	87.0	19.16
4237	0	0	0	269.0	133.5	83.0	21.47

	heartRate	glucose	TenYearCHD
0	80.0	77.0	0
1	95.0	76.0	0
2	75.0	70.0	0
3	65.0	103.0	1
4	85.0	85.0	0
...
4233	66.0	86.0	1
4234	65.0	68.0	0
4235	84.0	86.0	0
4236	86.0	NaN	0
4237	80.0	107.0	0

[4238 rows x 16 columns]

```
[13]: df.dtypes
```

```
[13]: male                int64
age                  int64
education            float64
currentSmoker        int64
cigsPerDay           float64
BPMeds               float64
prevalentStroke       int64
prevalentHyp          int64
diabetes              int64
totChol              float64
sysBP                float64
diaBP                float64
BMI                  float64
heartRate            float64
glucose              float64
TenYearCHD           int64
dtype: object
```

```
[14]: df=df.drop(["education"], axis = 1)
df
```

```
[14]:
```

	male	age	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	\
0	1	39	0	0.0	0.0	0	
1	0	46	0	0.0	0.0	0	

2	1	48		1	20.0	0.0		0
3	0	61		1	30.0	0.0		0
4	0	46		1	23.0	0.0		0
...
4233	1	50		1	1.0	0.0		0
4234	1	51		1	43.0	0.0		0
4235	0	48		1	20.0	NaN		0
4236	0	44		1	15.0	0.0		0
4237	0	52		0	0.0	0.0		0

	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	heartRate	\
0	0	0	195.0	106.0	70.0	26.97	80.0	
1	0	0	250.0	121.0	81.0	28.73	95.0	
2	0	0	245.0	127.5	80.0	25.34	75.0	
3	1	0	225.0	150.0	95.0	28.58	65.0	
4	0	0	285.0	130.0	84.0	23.10	85.0	
...
4233	1	0	313.0	179.0	92.0	25.97	66.0	
4234	0	0	207.0	126.5	80.0	19.71	65.0	
4235	0	0	248.0	131.0	72.0	22.00	84.0	
4236	0	0	210.0	126.5	87.0	19.16	86.0	
4237	0	0	269.0	133.5	83.0	21.47	80.0	

	glucose	TenYearCHD
0	77.0	0
1	76.0	0
2	70.0	0
3	103.0	1
4	85.0	0
...
4233	86.0	1
4234	68.0	0
4235	86.0	0
4236	NaN	0
4237	107.0	0

[4238 rows x 15 columns]

```
[15]: df.isna().any(axis=1).sum()
```

```
[15]: 489
```

```
[16]: print(f'{df.isna().any(axis=1).sum() / df.shape[0]:.2%}')
```

11.54%

there are 489 rows record has NaN value, and 11.54% of the total. We can just drop those reocrds.

```
[17]: df_new = df.dropna()
df_new
```

```
[17]:      male  age  currentSmoker  cigsPerDay  BPMeds  prevalentStroke  \
0         1   39                0          0.0    0.0                0
1         0   46                0          0.0    0.0                0
2         1   48                1         20.0    0.0                0
3         0   61                1         30.0    0.0                0
4         0   46                1         23.0    0.0                0
...
4231      1   58                0          0.0    0.0                0
4232      1   68                0          0.0    0.0                0
4233      1   50                1          1.0    0.0                0
4234      1   51                1         43.0    0.0                0
4237      0   52                0          0.0    0.0                0
```

```
      prevalentHyp  diabetes  totChol  sysBP  diaBP  BMI  heartRate  \
0                0         0   195.0  106.0   70.0  26.97      80.0
1                0         0   250.0  121.0   81.0  28.73      95.0
2                0         0   245.0  127.5   80.0  25.34      75.0
3                1         0   225.0  150.0   95.0  28.58      65.0
4                0         0   285.0  130.0   84.0  23.10      85.0
...
4231              1         0   187.0  141.0   81.0  24.96      80.0
4232              1         0   176.0  168.0   97.0  23.14      60.0
4233              1         0   313.0  179.0   92.0  25.97      66.0
4234              0         0   207.0  126.5   80.0  19.71      65.0
4237              0         0   269.0  133.5   83.0  21.47      80.0
```

```
      glucose  TenYearCHD
0         77.0          0
1         76.0          0
2         70.0          0
3        103.0          1
4         85.0          0
...
4231        81.0          0
4232        79.0          1
4233        86.0          1
4234        68.0          0
4237       107.0          0
```

[3749 rows x 15 columns]

```
[18]: features = df_new.iloc[:, :-1]
features
```

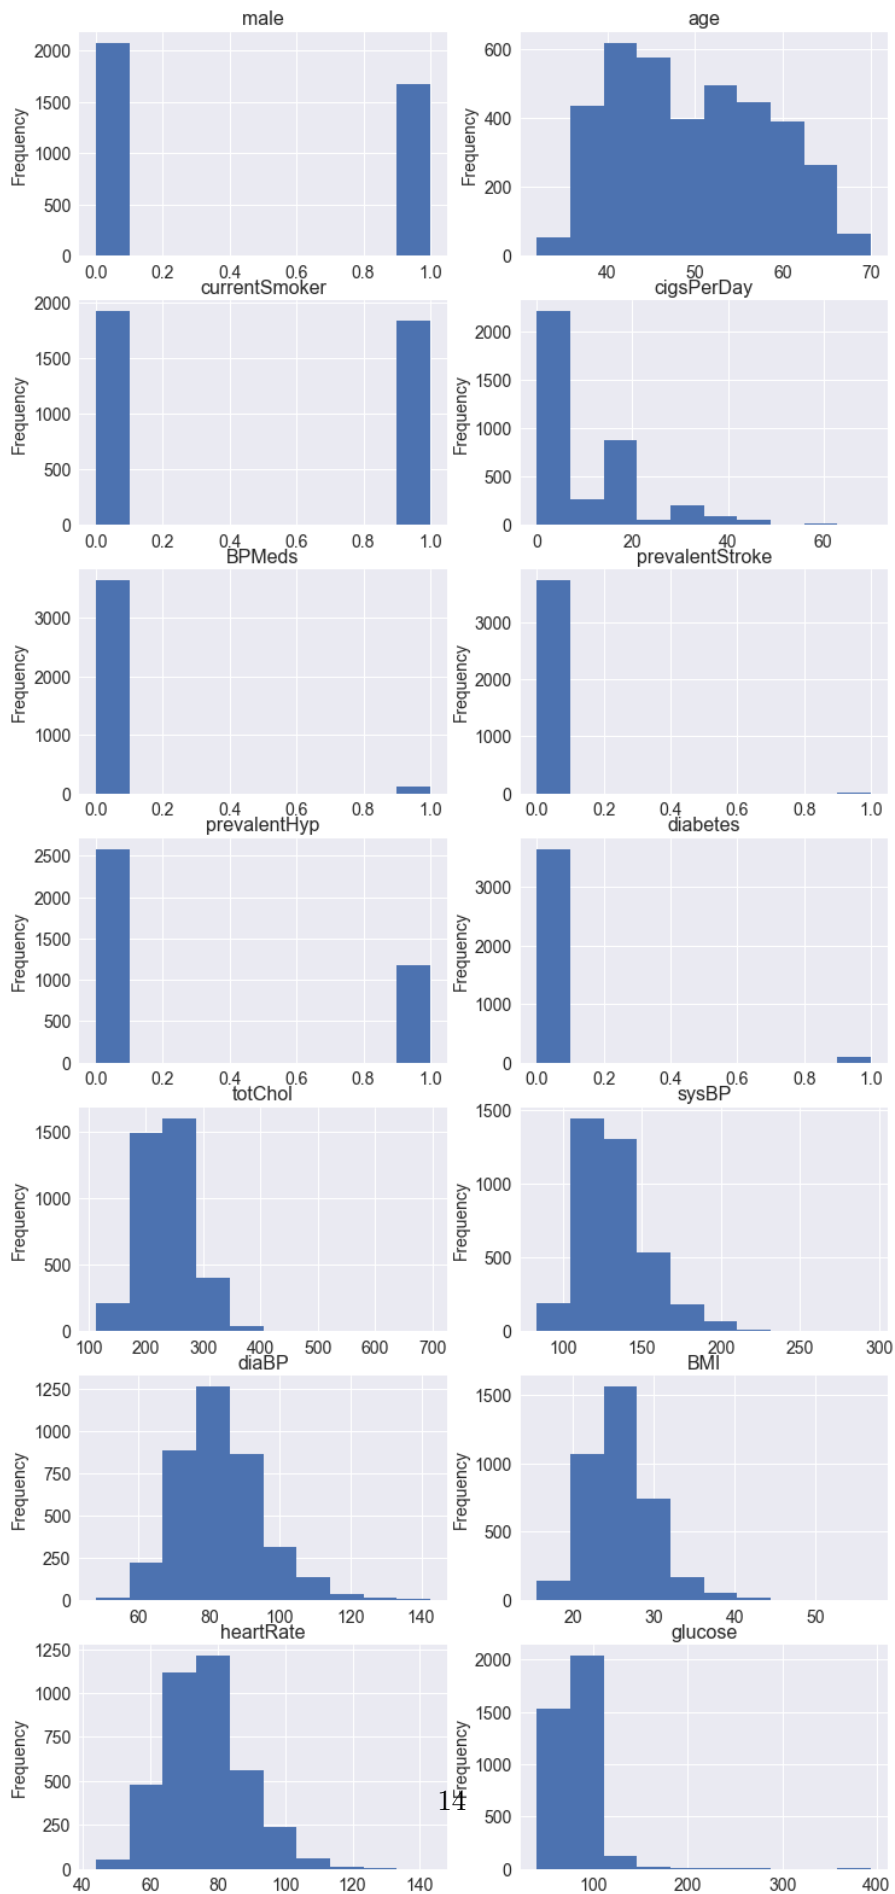
```
[18]:
```

	male	age	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	\
0	1	39	0	0.0	0.0	0	
1	0	46	0	0.0	0.0	0	
2	1	48	1	20.0	0.0	0	
3	0	61	1	30.0	0.0	0	
4	0	46	1	23.0	0.0	0	
...	
4231	1	58	0	0.0	0.0	0	
4232	1	68	0	0.0	0.0	0	
4233	1	50	1	1.0	0.0	0	
4234	1	51	1	43.0	0.0	0	
4237	0	52	0	0.0	0.0	0	

	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	heartRate	glucose
0	0	0	195.0	106.0	70.0	26.97	80.0	77.0
1	0	0	250.0	121.0	81.0	28.73	95.0	76.0
2	0	0	245.0	127.5	80.0	25.34	75.0	70.0
3	1	0	225.0	150.0	95.0	28.58	65.0	103.0
4	0	0	285.0	130.0	84.0	23.10	85.0	85.0
...
4231	1	0	187.0	141.0	81.0	24.96	80.0	81.0
4232	1	0	176.0	168.0	97.0	23.14	60.0	79.0
4233	1	0	313.0	179.0	92.0	25.97	66.0	86.0
4234	0	0	207.0	126.5	80.0	19.71	65.0	68.0
4237	0	0	269.0	133.5	83.0	21.47	80.0	107.0

[3749 rows x 14 columns]

```
[19]: fig, axes= plt.subplots(7,2,figsize=(12,7*4))
flat_axes = [item for sub in axes for item in sub]
iter_ax = iter(flat_axes)
for f_col in features.columns:
    ax=next(iter_ax)
    features[f_col].plot(kind='hist', ax=ax)
    ax.set_title(f_col)
plt.show()
```



```
[20]: y = df_new.iloc[:, -1]
      y
```

```
[20]: 0      0
      1      0
      2      0
      3      1
      4      0
      ..
     4231     0
     4232     1
     4233     1
     4234     0
     4237     0
      Name: TenYearCHD, Length: 3749, dtype: int64
```

```
[21]: clf = LogisticRegression(solver='newton-cg')
      clf.fit(features, y)
      print(clf.coef_)
      print(clf.intercept_)

[[ 0.5674254  0.0640999  0.0710263  0.01848422  0.14415091  0.58079615
    0.2122989  0.00221661  0.00223462  0.01534348 -0.00382703  0.0103664
   -0.00234315  0.00758097]]
[-8.64869189]
```

```
[22]: clf2 = LogisticRegression(solver='liblinear')
      clf2.fit(features, y)
      print(clf.coef_)
      print(clf.intercept_)

[[ 0.5674254  0.0640999  0.0710263  0.01848422  0.14415091  0.58079615
    0.2122989  0.00221661  0.00223462  0.01534348 -0.00382703  0.0103664
   -0.00234315  0.00758097]]
[-8.64869189]
```

```
[23]: st_x = sm.add_constant(features)
      log_reg = sm.Logit(y, st_x)
      log_result = log_reg.fit()
      log_result.summary()
      # for the warning of Pandas, we can change the data from dataframe and series to
      ↪ numpy ndarray to solve it. but it did not make difference to the results.
```

```
Optimization terminated successfully.
      Current function value: 0.377199
      Iterations 7
```

C:\Python39\lib\site-packages\statsmodels\tsa\tsatools.py:142: FutureWarning: In a future version of pandas all arguments of concat except for the argument 'objs' will be keyword-only

```
x = pd.concat(x[:, :order], 1)
```

[23]: <class 'statsmodels.iolib.summary.Summary'>

```

"""
                                Logit Regression Results
=====
Dep. Variable:                TenYearCHD    No. Observations:                3749
Model:                        Logit         Df Residuals:                  3734
Method:                       MLE          Df Model:                   14
Date:                         Tue, 28 Sep 2021    Pseudo R-squ.:                0.1169
Time:                         09:59:57         Log-Likelihood:               -1414.1
converged:                    True            LL-Null:                     -1601.4
Covariance Type:              nonrobust        LLR p-value:                  2.922e-71
=====
===
                                coef      std err          z      P>|z|      [0.025
0.975]
-----
---
const                -8.6463      0.687     -12.577      0.000     -9.994
-7.299
male                  0.5740      0.107      5.343      0.000      0.363
0.785
age                   0.0640      0.007      9.787      0.000      0.051
0.077
currentSmoker         0.0732      0.155      0.473      0.636     -0.230
0.376
cigsPerDay            0.0184      0.006      3.003      0.003      0.006
0.030
BPMeds                0.1446      0.232      0.622      0.534     -0.311
0.600
prevalentStroke        0.7191      0.489      1.471      0.141     -0.239
1.677
prevalentHyp           0.2146      0.136      1.574      0.116     -0.053
0.482
diabetes              0.0025      0.312      0.008      0.994     -0.609
0.614
totChol               0.0022      0.001      2.074      0.038      0.000
0.004
sysBP                 0.0153      0.004      4.080      0.000      0.008
0.023
diaBP                -0.0039      0.006     -0.619      0.536     -0.016
0.009
BMI                   0.0103      0.013      0.820      0.412     -0.014

```



```

0.035
heartRate      -0.0023      0.004      -0.550      0.583      -0.010
0.006
glucose        0.0076      0.002      3.408      0.001      0.003
0.012
=====
===
"""

```

for many variables, the cooresponding P-Value is high, thus we should accept the hypothesis that “the related cooefficient is null/zero”

Those we should keep are:

- gender
- age
- how many cigs per day
- cholesterol
- systolic blood pressure
- glucose

Those we may consider are:

- prevalent Hypertensive
- prevalent stroke

Thus we can keep only those 8 features to do the regression again.

```

[24]: feature_cols = ['male', 'age', 'cigsPerDay', 'prevalentStroke', 'prevalentHyp', '
      ↪ 'totChol', 'sysBP', 'glucose']
features2 = df_new[feature_cols]

clf = LogisticRegression(solver='newton-cg')
clf.fit(features2, y)
print(clf.coef_)
print(clf.intercept_)

```

```

[[0.56711658 0.06490845 0.02006905 0.61631235 0.21189139 0.00222781
  0.01427033 0.00762089]]
[-8.74987856]

```

```

[25]: y_predict = clf.predict(features2)
check = y==y_predict
check

```

```

[25]: 0      True
      1      True
      2      True
      3     False
      4      True
      ...

```

```

4231      True
4232     False
4233     False
4234      True
4237      True
Name: TenYearCHD, Length: 3749, dtype: bool

```

```
[26]: check.sum(), check.sum()/check.shape[0]
```

```
[26]: (3203, 0.8543611629767938)
```

Out of total 3749 records, the module build only correct for 3203, 85.44% of total

```

[27]: st_x2 = sm.add_constant(features2)
log_reg = sm.Logit(y, st_x2)
log_result = log_reg.fit()
log_result.summary()
# for the warning of Pandas, we can change the data from dataframe and series to
→numpy ndarray to solve it. but it did not make difference to the results.

```

Optimization terminated successfully.

Current function value: 0.377440

Iterations 7

C:\Python39\lib\site-packages\statsmodels\tsa\tsatools.py:142: FutureWarning: In a future version of pandas all arguments of concat except for the argument

'objs' will be keyword-only

```
x = pd.concat(x[:,order], 1)
```

```
[27]: <class 'statsmodels.iolib.summary.Summary'>
```

```

"""
                                Logit Regression Results
=====
Dep. Variable:                  TenYearCHD    No. Observations:                  3749
Model:                            Logit      Df Residuals:                  3740
Method:                           MLE       Df Model:                        8
Date:                Tue, 28 Sep 2021    Pseudo R-squ.:                  0.1164
Time:                   09:59:58    Log-Likelihood:                 -1415.0
converged:                      True     LL-Null:                       -1601.4
Covariance Type:            nonrobust    LLR p-value:                   1.280e-75
=====
===
                                coef    std err          z      P>|z|      [0.025
0.975]
-----
---
const                -8.7497      0.515    -16.986      0.000     -9.759
-7.740
male                  0.5732      0.105     5.434      0.000      0.366

```

0.780					
age	0.0649	0.006	10.222	0.000	0.052
0.077					
cigsPerDay	0.0200	0.004	4.877	0.000	0.012
0.028					
prevalentStroke	0.7615	0.483	1.575	0.115	-0.186
1.709					
prevalentHyp	0.2138	0.133	1.601	0.109	-0.048
0.475					
totChol	0.0022	0.001	2.076	0.038	0.000
0.004					
sysBP	0.0142	0.003	5.022	0.000	0.009
0.020					
glucose	0.0076	0.002	4.603	0.000	0.004
0.011					

```

=====
===
"""

```

As statsmodels prediction is continuous variable not good to do the check here as our target is discrete.

```
[28]: y_predict_sm = log_result.predict(st_x2)
      y_predict_sm
```

```
[28]: 0      0.042487
      1      0.051892
      2      0.146024
      3      0.364696
      4      0.102520
      ...
      4231   0.238984
      4232   0.458785
      4233   0.310810
      4234   0.227012
      4237   0.113098
      Length: 3749, dtype: float64
```

```
[ ]:
```